

Trees & Forests

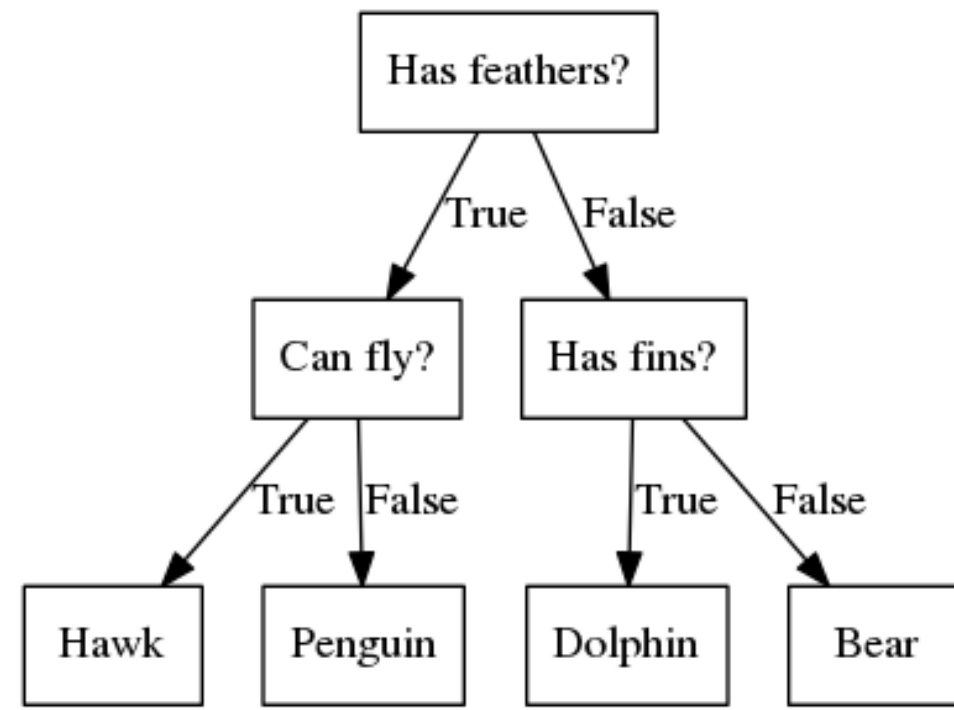
Andreas Müller

Why trees?

- Very powerful modeling method – non-linear!
- Doesn't care about scaling of distribution of data!
- “Interpretable”
- Basis of very powerful models!

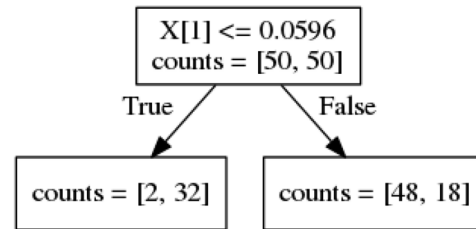
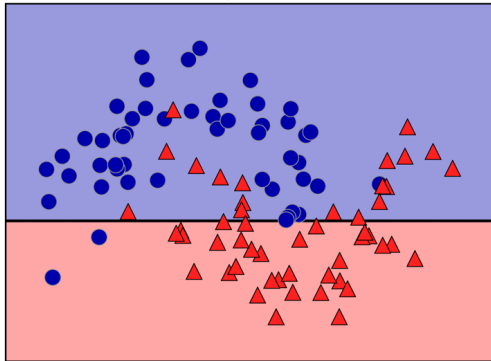
Decision Trees for Classification

Idea: series of binary questions

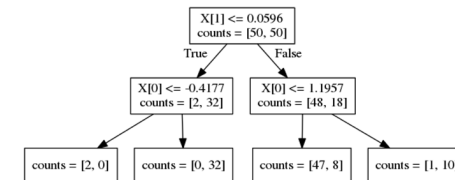
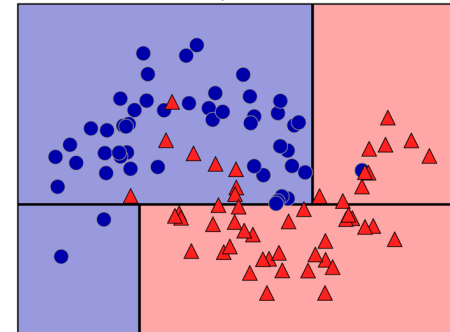


Building trees

depth = 1



depth = 2



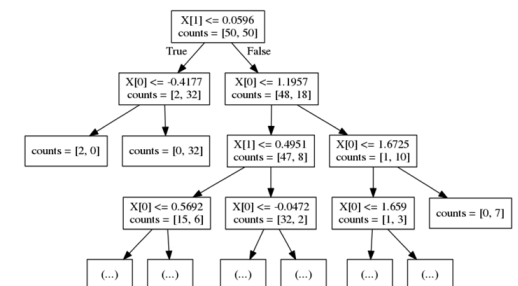
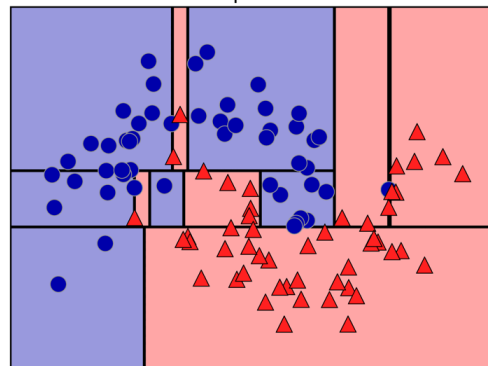
Continuous features:
“questions” are thresholds on
single features.

[Other methods are possible
but not as common]

For each split:
exhaustive search over all
features and thresholds!

Minimize “impurity”

depth = 9



Criteria (for classification)

- Gini index:

$$H_{\text{gini}}(X_m) = \sum_{k \in \mathcal{Y}} p_{mk} (1 - p_{mk})$$

- Cross-entropy:

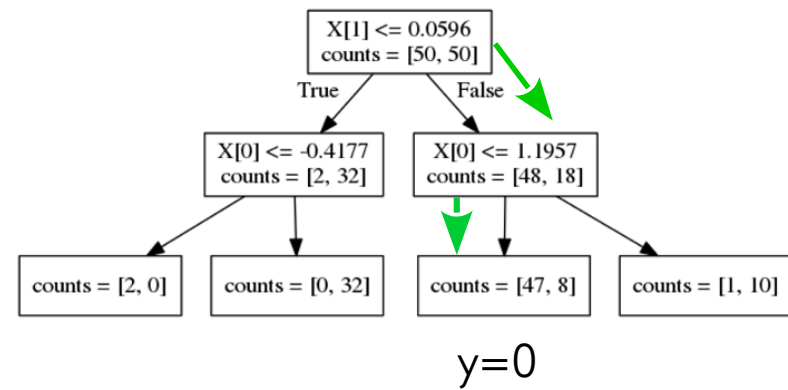
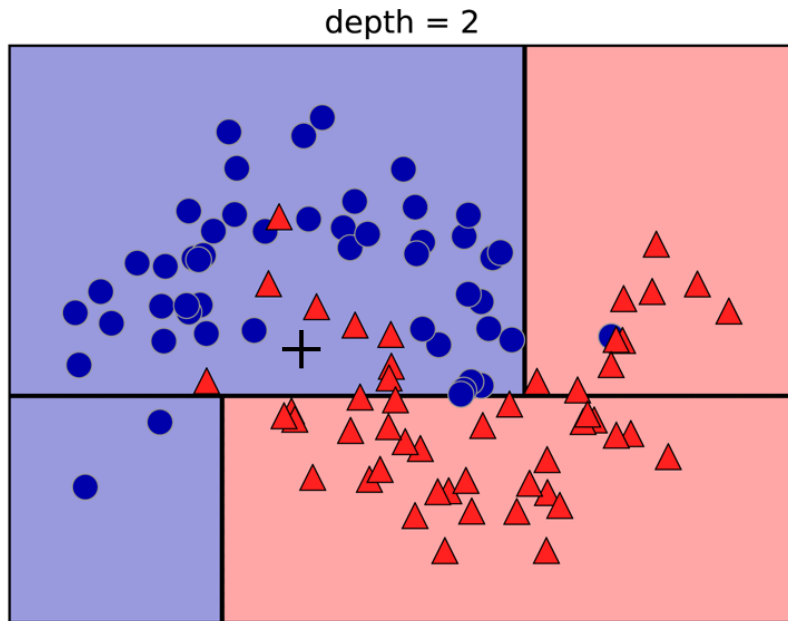
$$H_{\text{CE}}(X_m) = - \sum_{k \in \mathcal{Y}} p_{mk} \log(p_{mk})$$

X_m observations in node m

\mathcal{Y} classes

$p_{m\cdot}$ distribution over classes in node m

Prediction



Traverse tree based on feature tests
Predict most common class in leaf

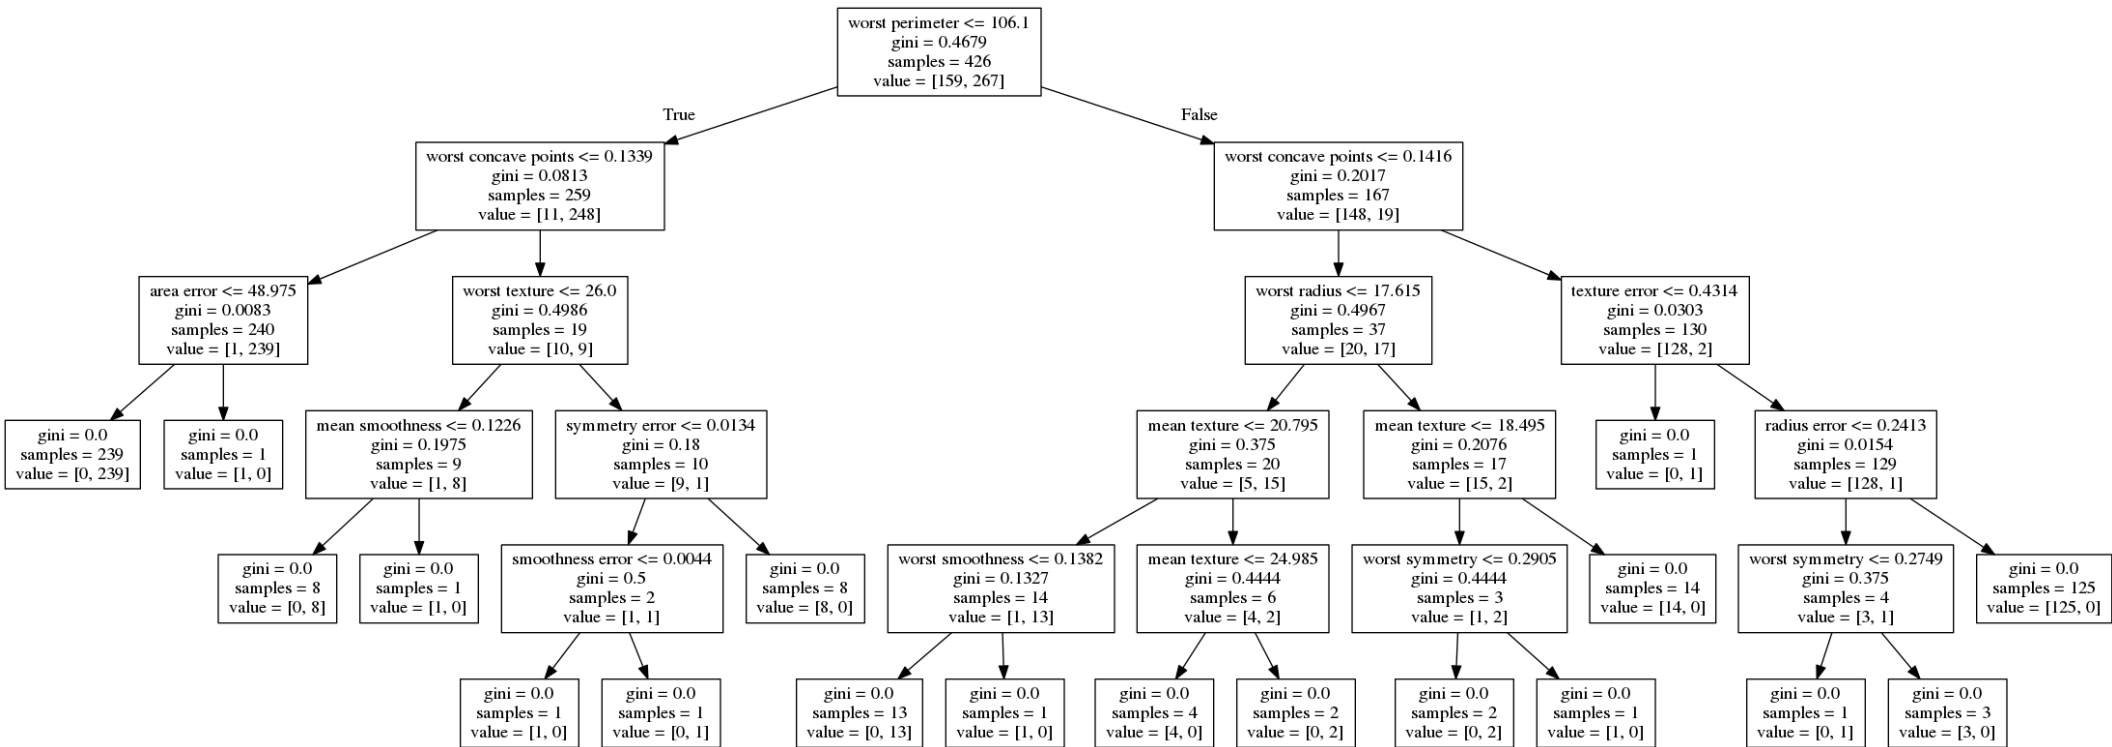
Regression trees

- Impurity Criteria:
Mean Squared Error
Mean Absolute Error
- Prediction:
Predict mean.
- Without regularization / pruning:
Each leaf often contains a single point to be “pure”

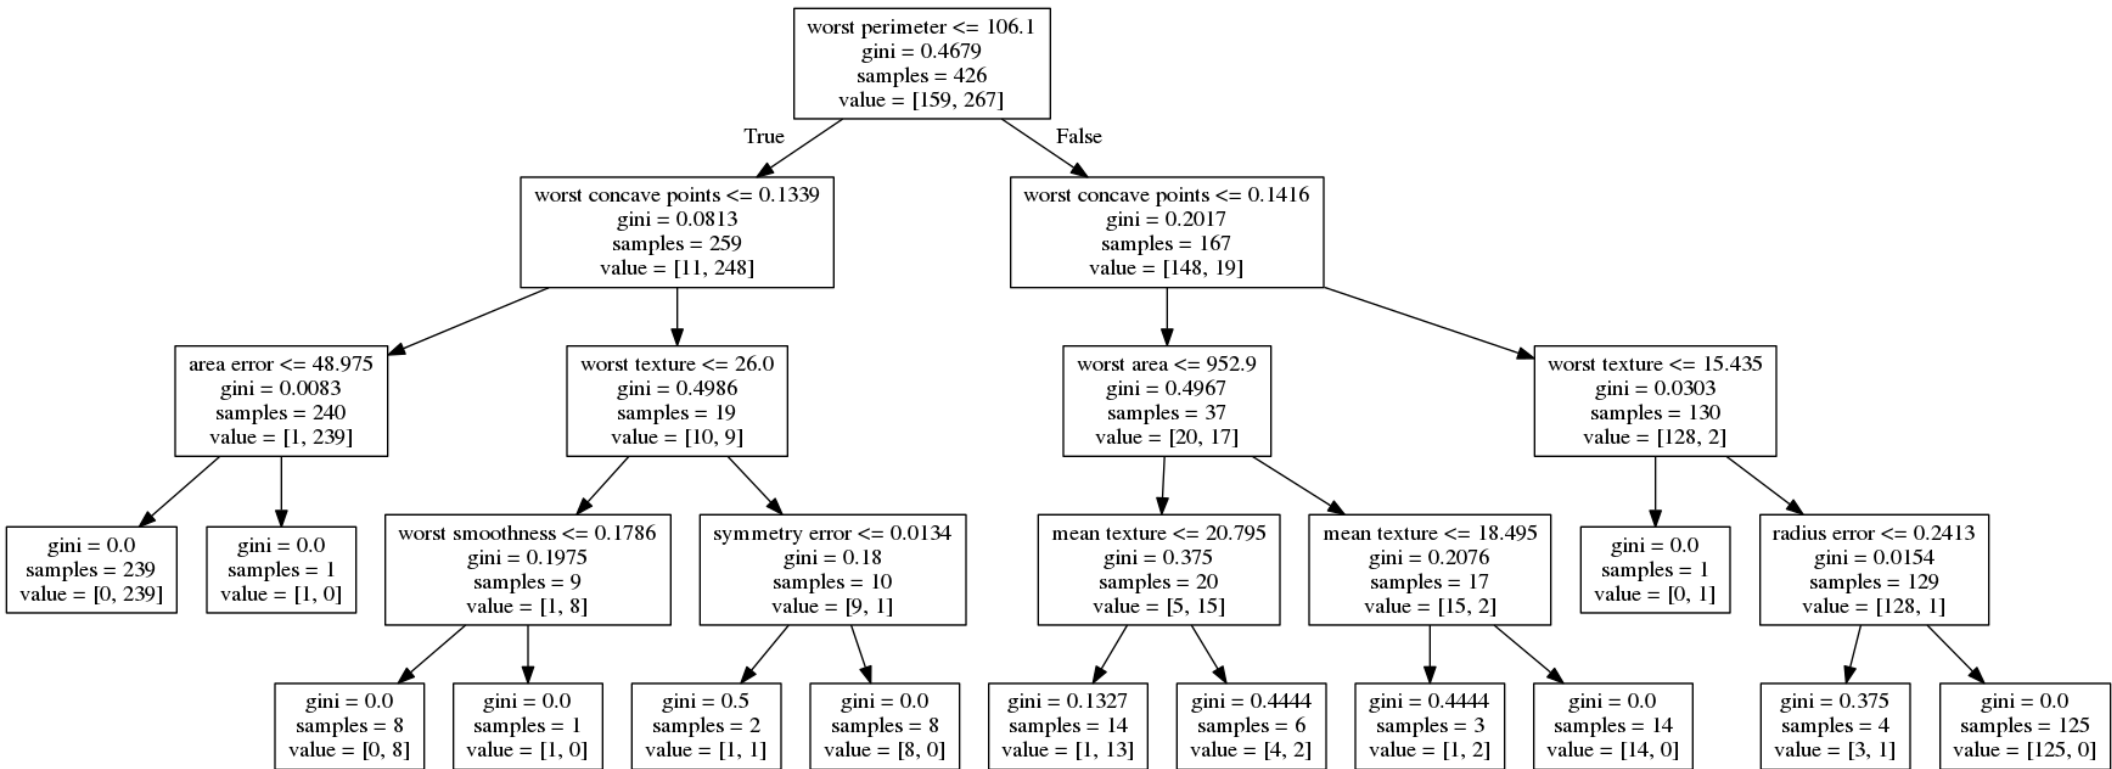
Parameter tuning

- Pre-pruning and post-pruning (not in sklearn yet)
- Limit tree size (pick one):
 - max_depth
 - max_leaf_nodes
 - min_samples_split
 - (and more)

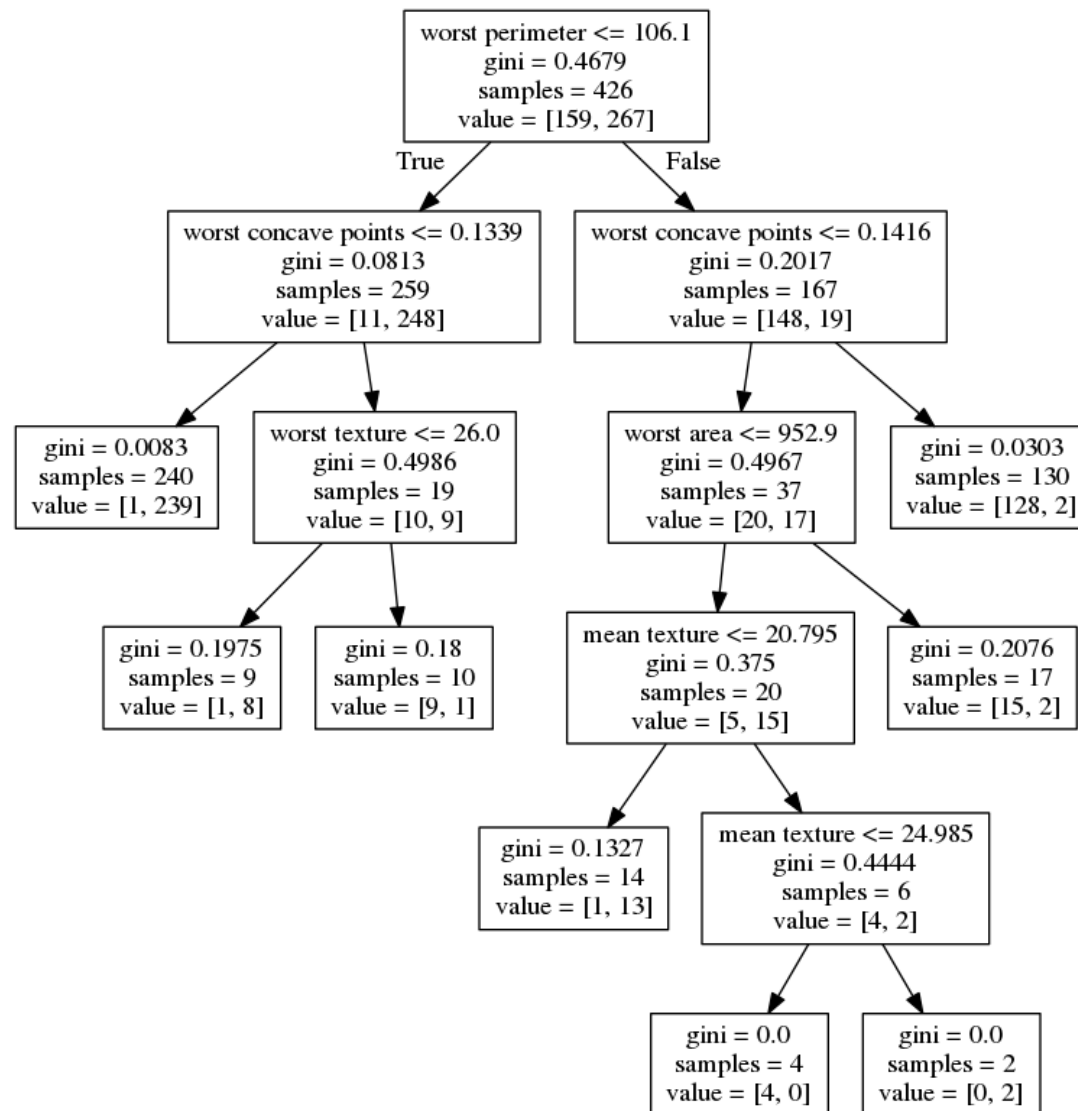
No pruning



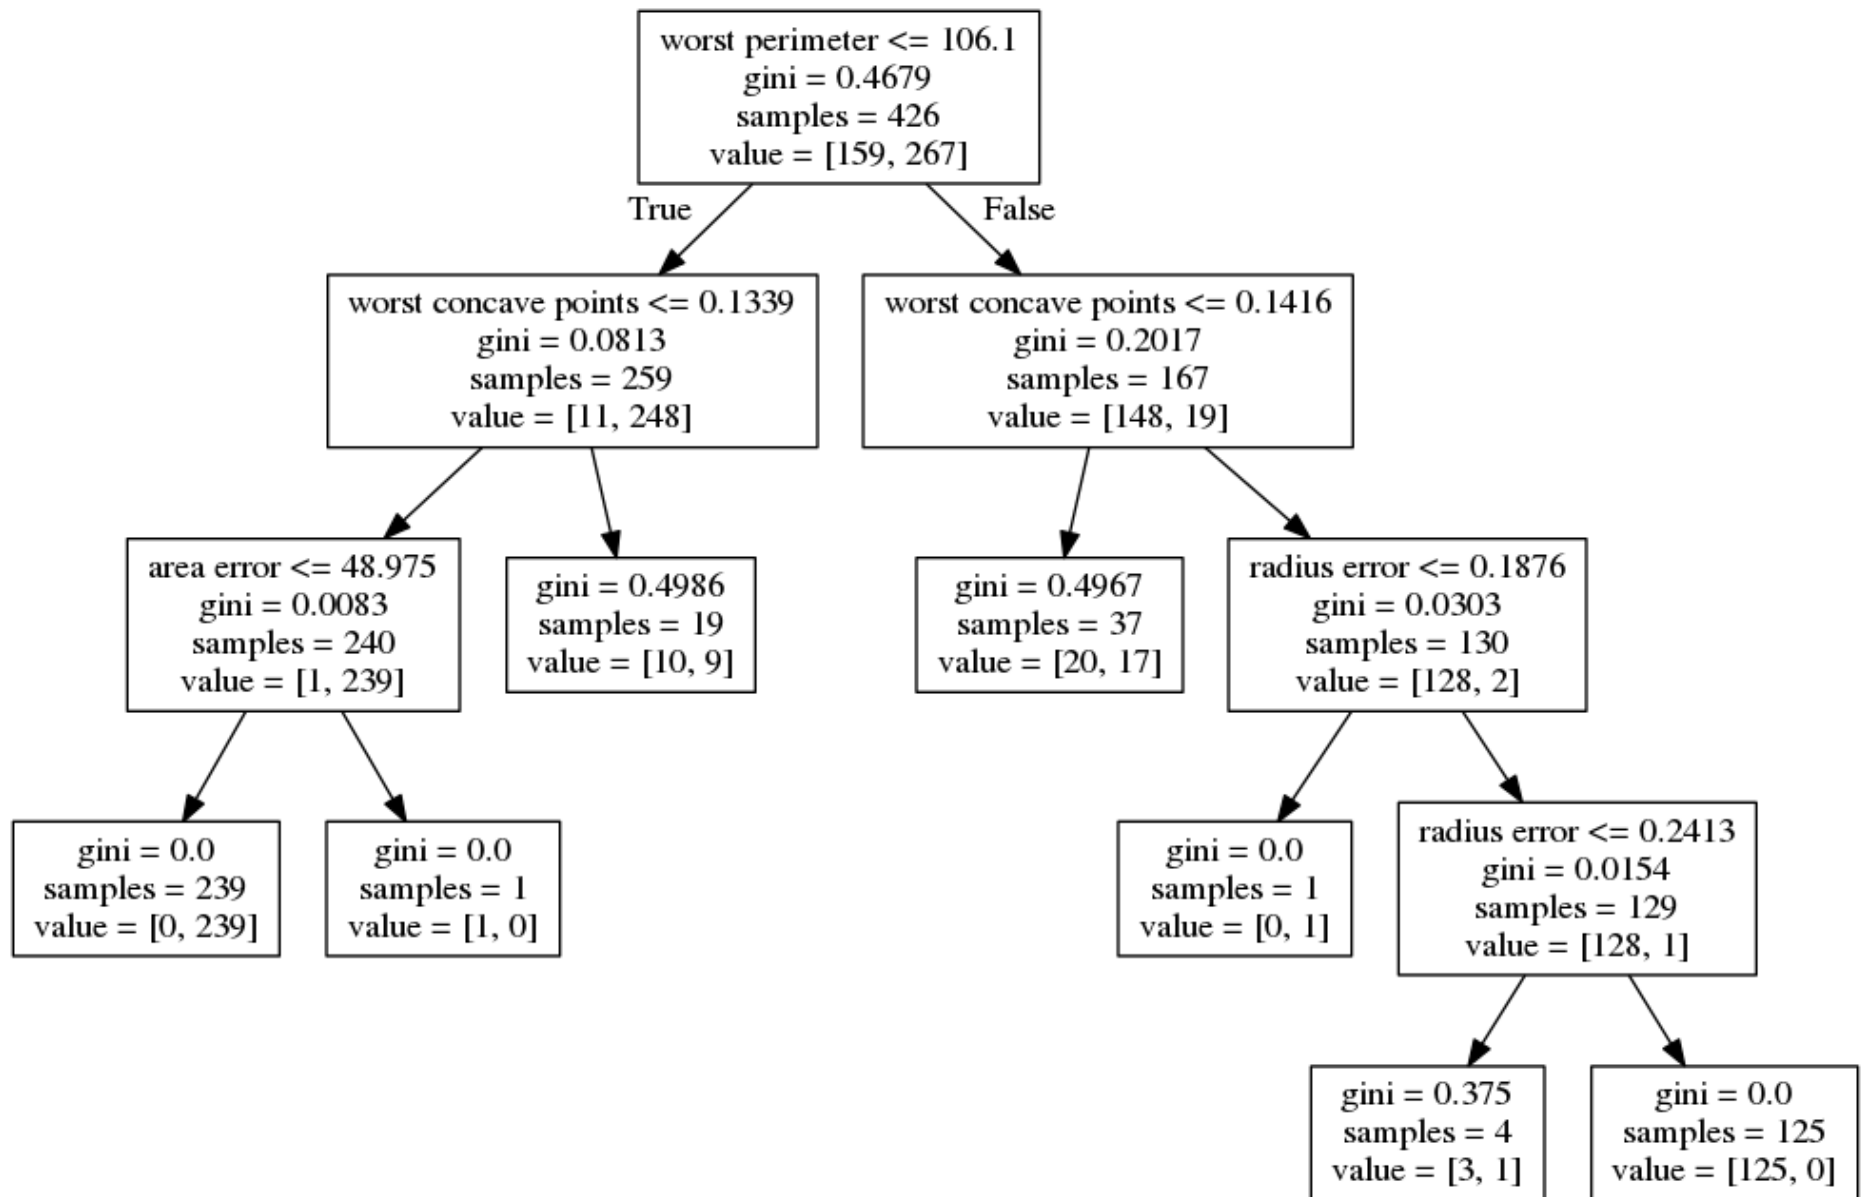
max_depth=4



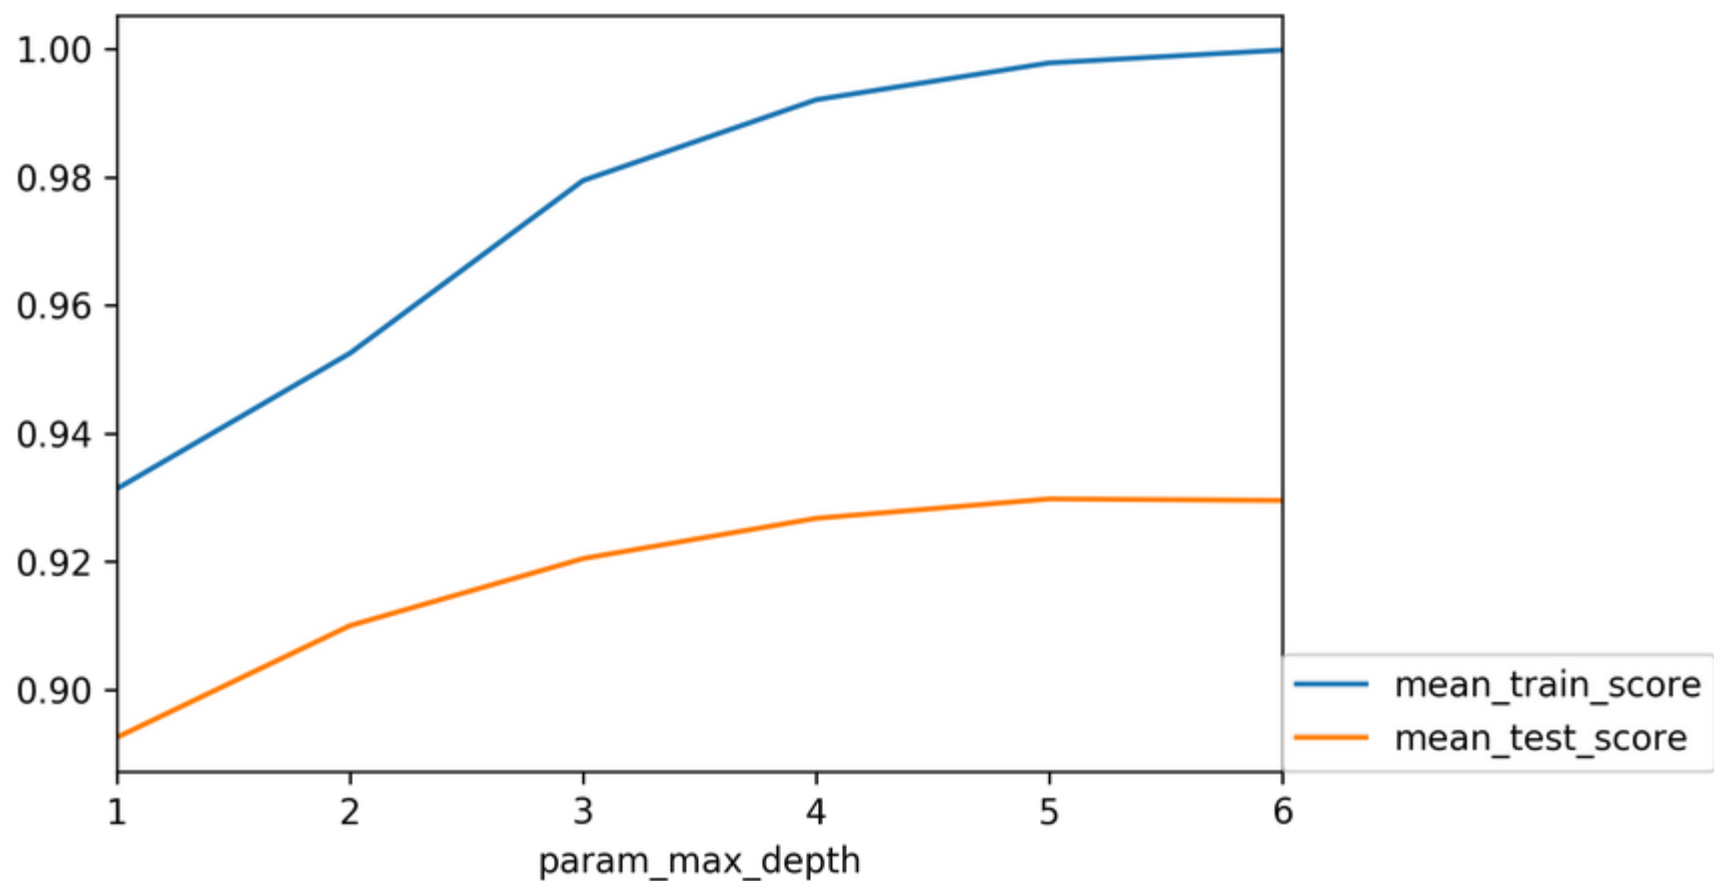
max_leaf_nodes=8



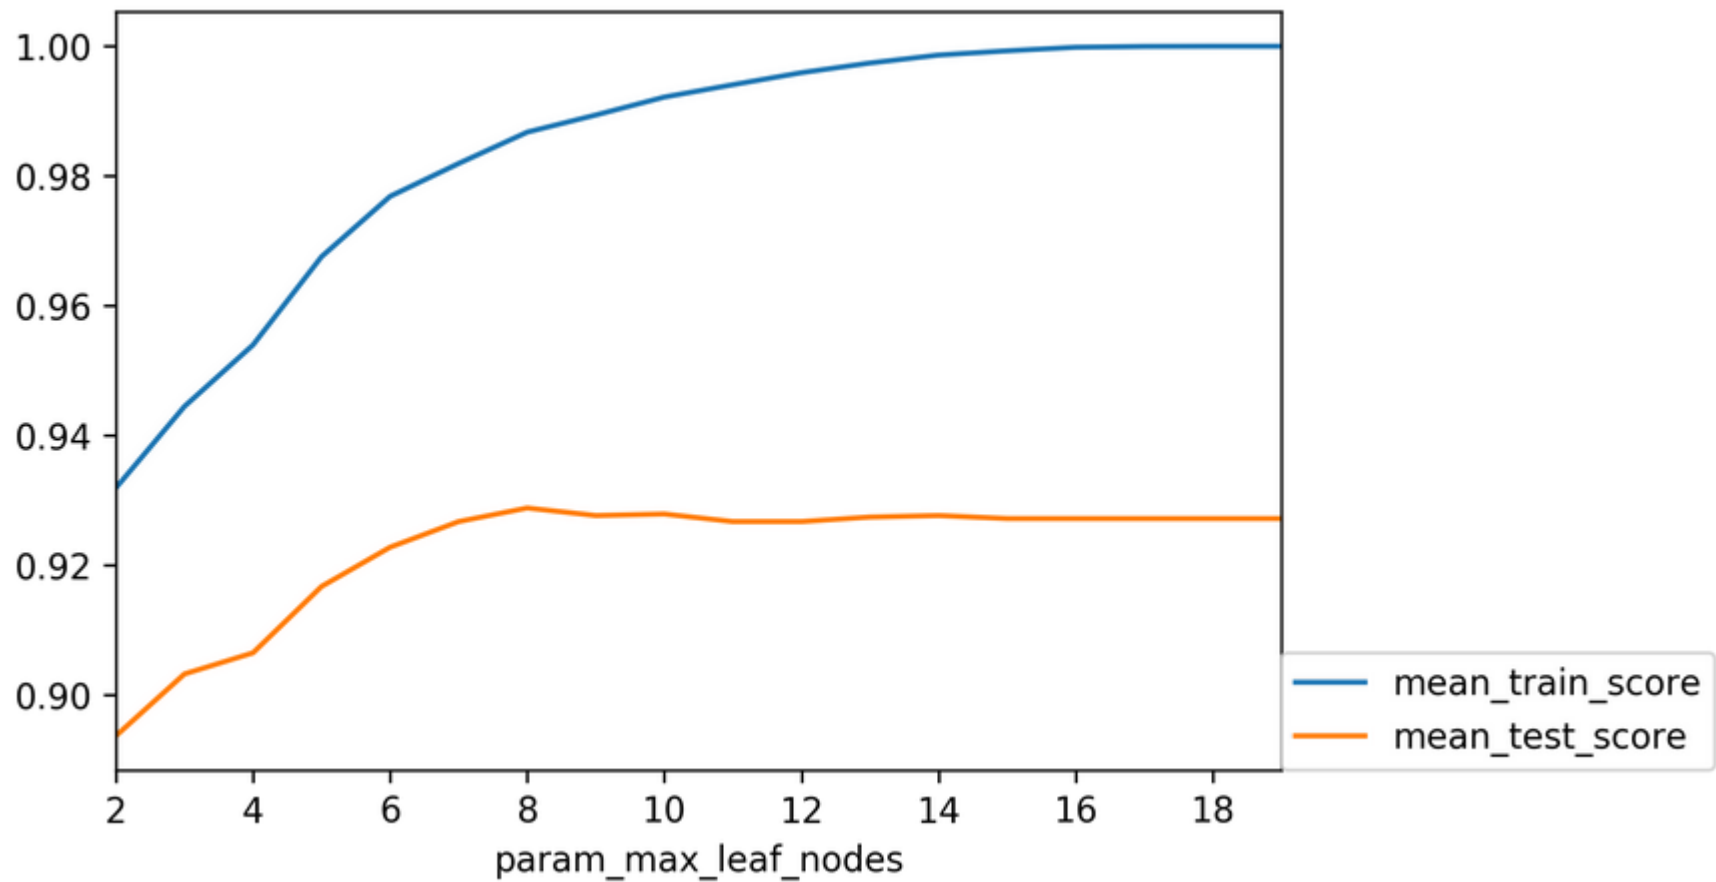
min_samples_split=50

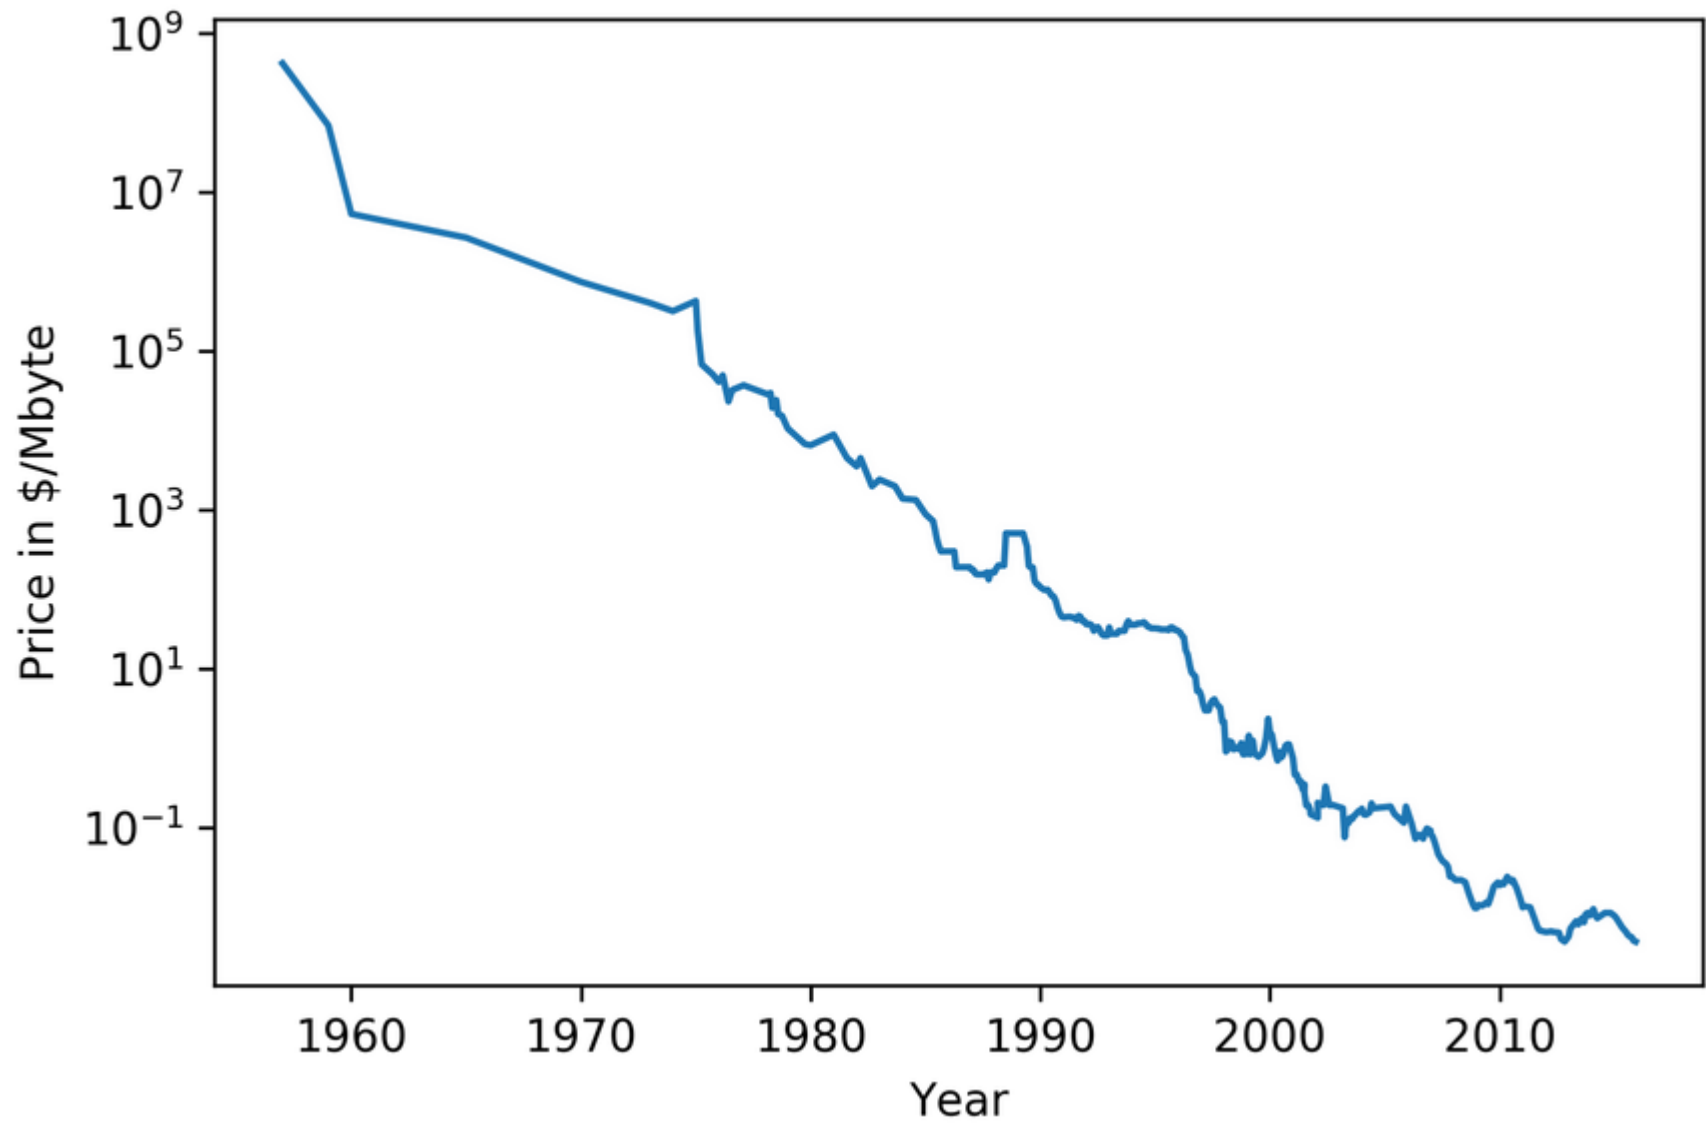


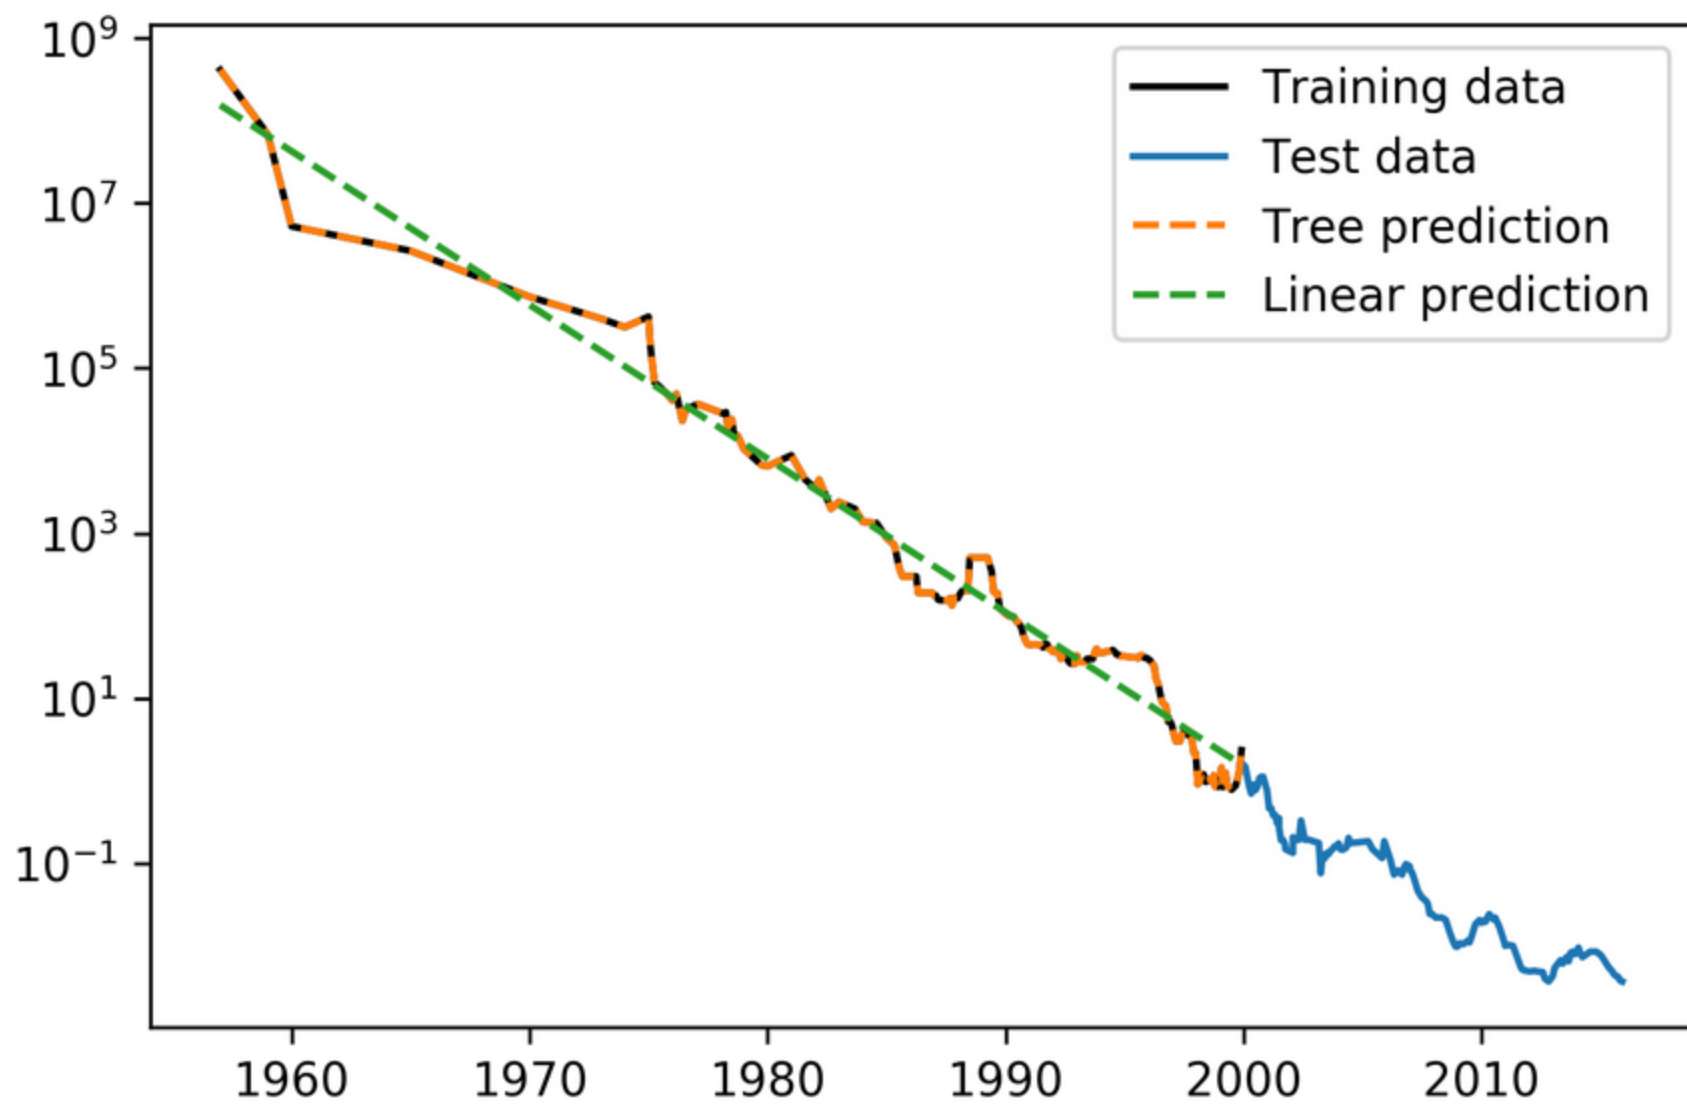
```
from sklearn.model_selection import GridSearchCV
param_grid = {'max_depth': range(1, 7)}
grid = GridSearchCV(DecisionTreeClassifier(random_state=0), param_grid=param_grid, cv=10)
grid.fit(X_train, y_train)
```



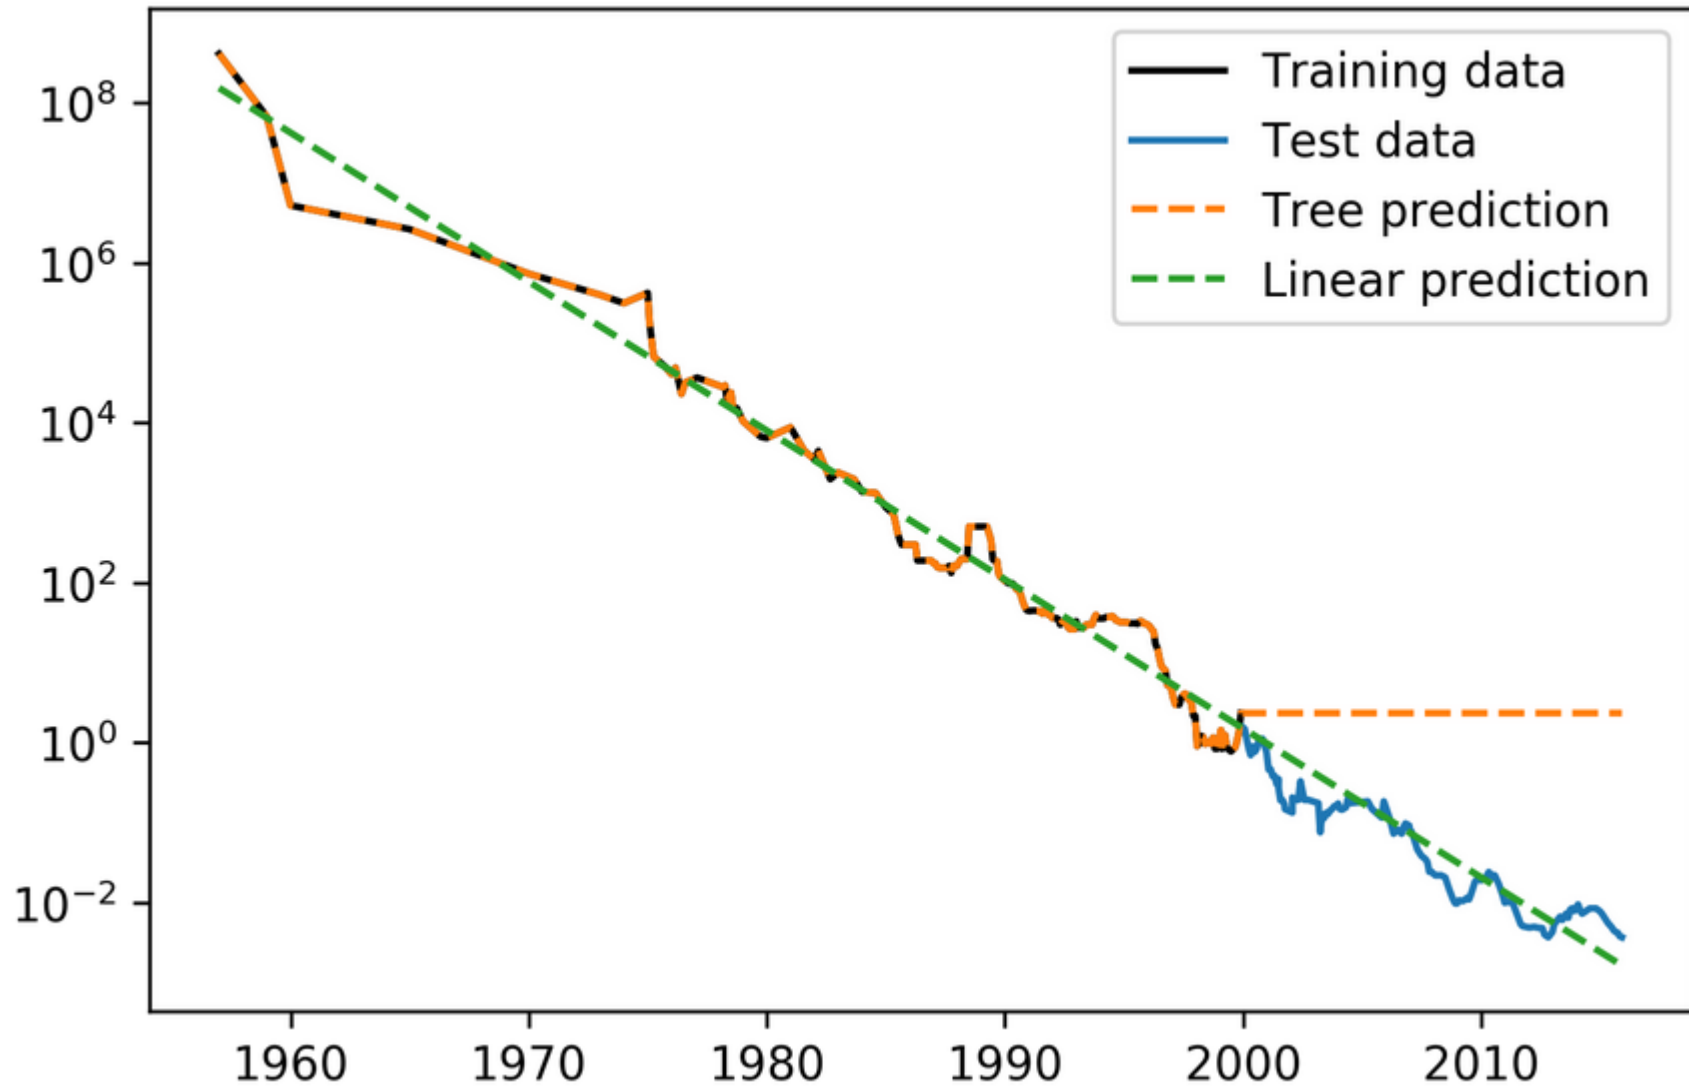
```
from sklearn.model_selection import GridSearchCV
param_grid = {'max_leaf_nodes':range(2, 20)}
grid = GridSearchCV(DecisionTreeClassifier(random_state=0), param_grid=param_grid, cv=10)
grid.fit(X_train, y_train)
```







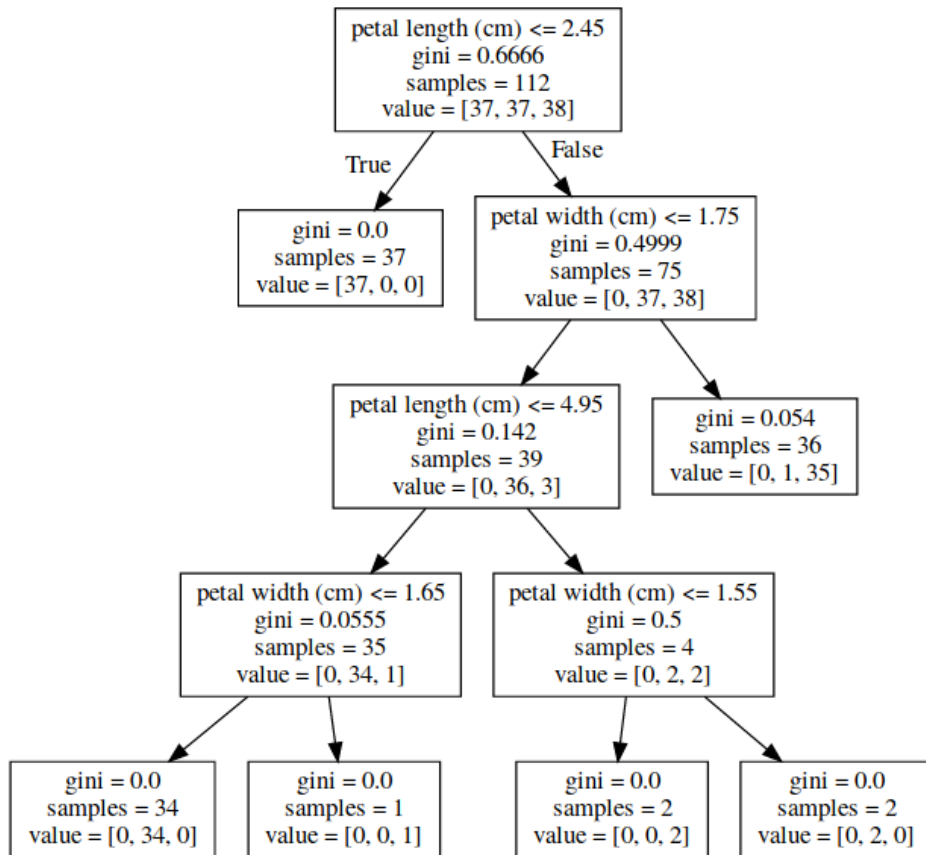
Extrapolation



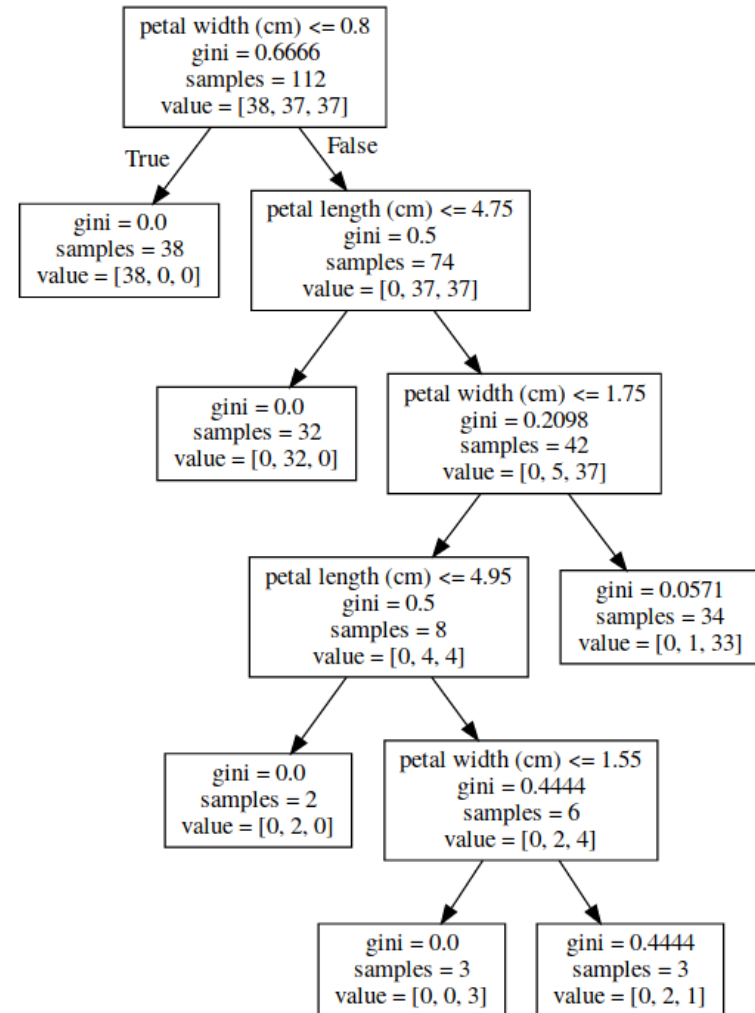
Would be the same for nearest neighbor regression!

Instability

```
iris = load_iris()
X_train, X_test, y_train, y_test = train_test_split(
    iris.data, iris.target, stratify=iris.target, random_state=0)
tree = DecisionTreeClassifier(max_leaf_nodes=6).fit(X_train, y_train)
tree_dot = export_graphviz(tree, out_file=None, feature_names=iris.feature_names)
graphviz.Source(tree_dot)
```

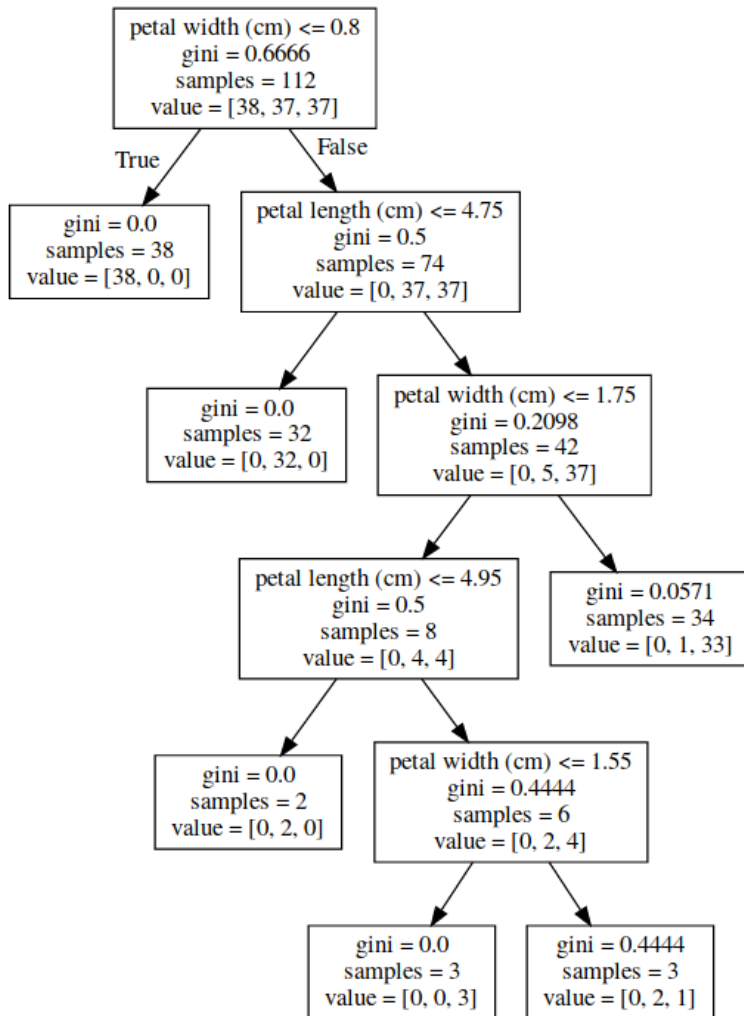


```
X_train, X_test, y_train, y_test = train_test_split(
    iris.data, iris.target, stratify=iris.target, random_state=1)
tree = DecisionTreeClassifier(max_leaf_nodes=6).fit(X_train, y_train)
tree_dot = export_graphviz(tree, out_file=None, feature_names=iris.feature_names)
graphviz.Source(tree_dot)
```



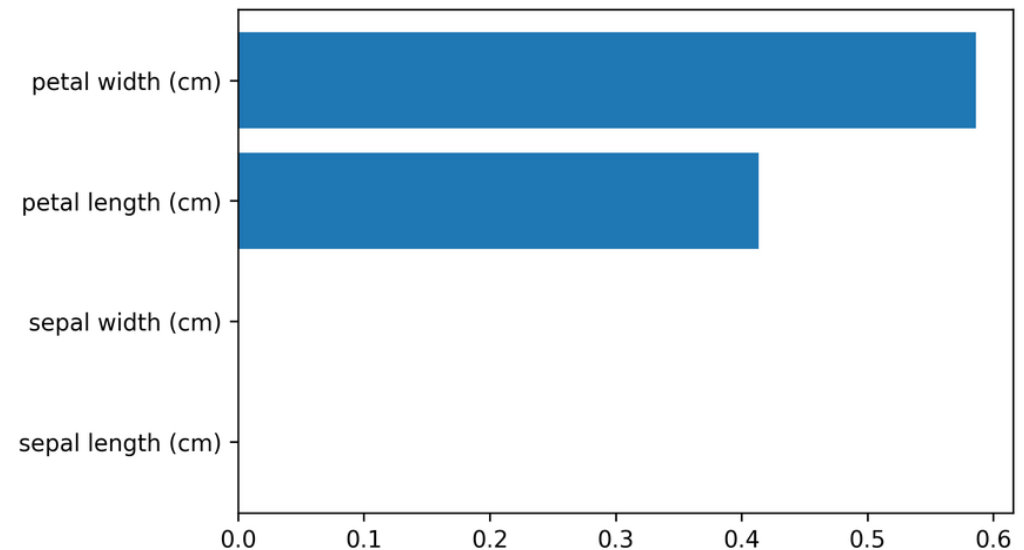
Feature importance

```
X_train, X_test, y_train, y_test = train_test_split(
    iris.data, iris.target, stratify=iris.target, random_state=1)
tree = DecisionTreeClassifier(max_leaf_nodes=6).fit(X_train, y_train)
tree_dot = export_graphviz(tree, out_file=None, feature_names=iris.feature_names)
graphviz.Source(tree_dot)
```



```
tree.feature_importances_
```

```
array([ 0.    ,  0.    ,  0.414,  0.586])
```



Unstable tree → unstable feature importances.

Might take one or multiple from a group of correlated features.

Categorical Data

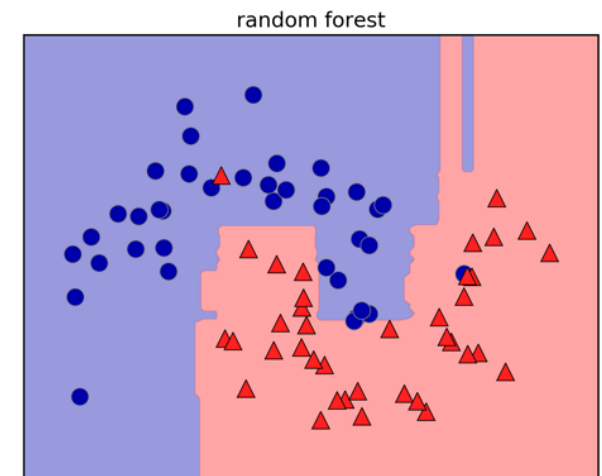
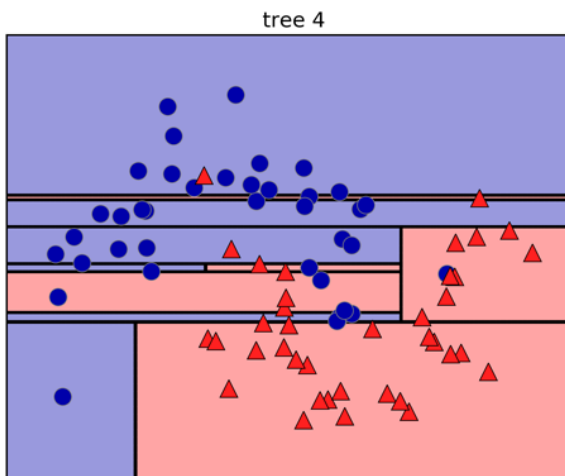
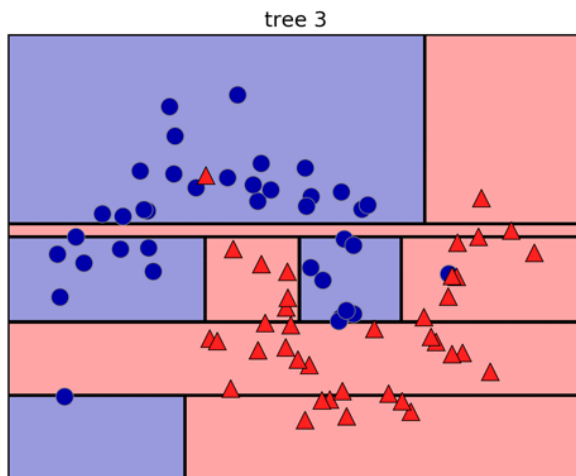
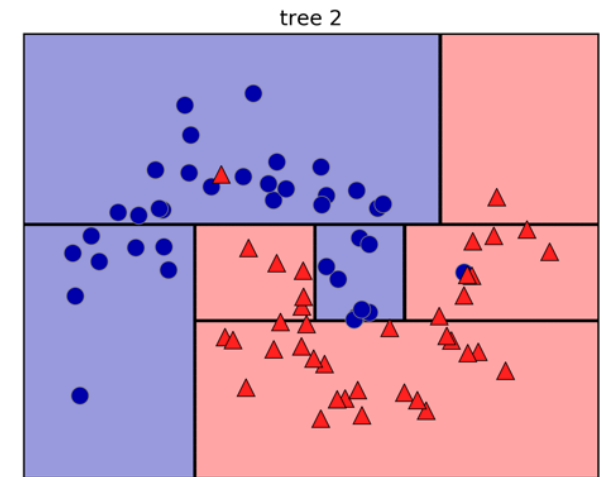
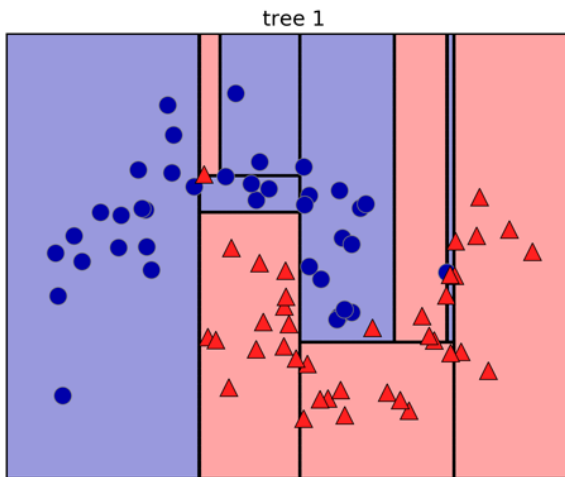
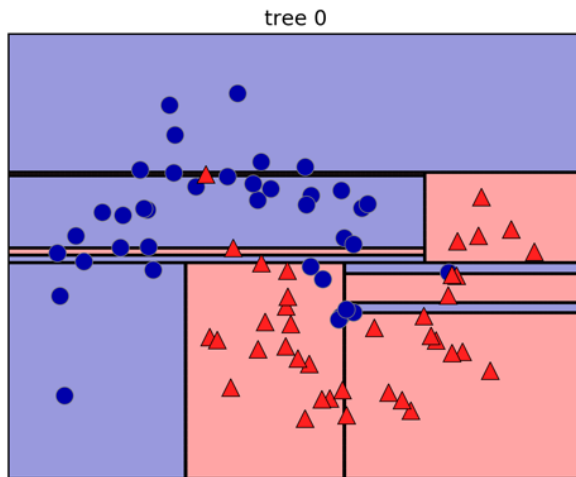
- Can split on categorical data directly
- Intuitive way to split: split in two subsets
- 2^{n_values} many possibilities
- Possible to do in linear time exactly for gini index and binary classification.
- Heuristics done in practice for multi-class.
- Not in sklearn release version :(

Predicting probabilities

- Fraction of class in leaf.
- Without pruning: Always 100% certain!
- Even with pruning might be too certain.

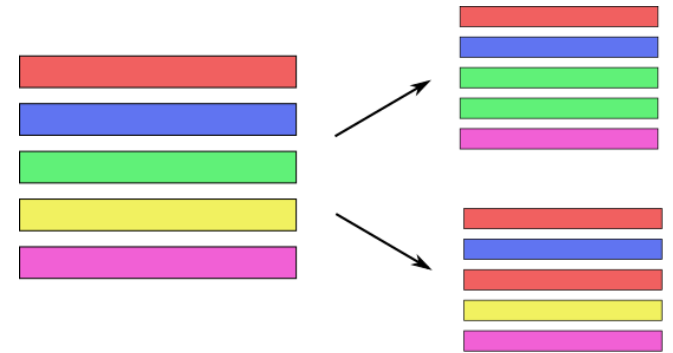
Random Forests

- Smarter bagging for trees!

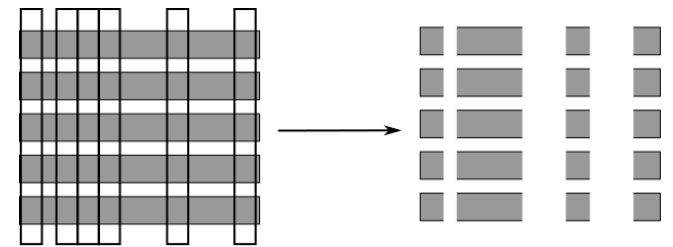


Randomize in two ways

- For each **tree**:
Pick bootstrap sample of data



- For each **split**:
Pick random sample of features



- More tree are always better

Tuning Random Forests

- Main parameter: `max_features`
 - around $\sqrt{n_features}$ for classification
 - Around `n_features` for regression
- `n_estimators > 100`
- Prepruning might help, definitely helps with model size!
- `max_depth`, `max_leaf_nodes`, `min_samples_split` again

Variable Importance

```
X_train, X_test, y_train, y_test = train_test_split(  
    iris.data, iris.target, stratify=iris.target, random_state=1)  
rf = RandomForestClassifier(n_estimators=100).fit(X_train, y_train)
```

```
rf.feature_importances_  
array([ 0.101,  0.034,  0.437,  0.428])
```

```
plt.barh(range(4), rf.feature_importances_)  
plt.yticks(range(4), iris.feature_names);
```

